

Soban Shankar

Chennai, India

+91 9444389781 | ssoban2004@gmail.com | [Portfolio](#) | [GitHub](#) | [LinkedIn](#)

PROFESSIONAL SUMMARY

AI Engineer with hands-on enterprise AI experience through a 6-month internship and multiple end-to-end Generative AI projects. Skilled in building AI agents, RAG systems, and evaluation pipelines, including browser-automation agents and multimodal (vision, speech, OCR) AI. Proficient in Python, LangChain, LangGraph, and vector databases.

EXPERIENCE

AI Engineering Intern | Drivestream | Chennai, India

Jan 2026 – Jul 2026

- Designed a multi-agent orchestration system to automate enterprise Oracle Fusion cloud transformation workflows, reducing manual configuration effort by approximately 60–70%.
- Built an AI browser agent that reads and validates Excel configuration data (combined code and LLM checks), then reasons over the live Oracle Fusion UI to decide each navigation step and apply changes autonomously across 20+ HCM and ERP modules, with retry logic and human-in-the-loop escalation for edge cases; engineered the agent's prompt logic.
- Built a conversational AI report agent using Retrieval-Augmented Generation (RAG) over 2,000+ Oracle Fusion HCM table and view definitions, enabling schema-aware SQL generation and natural language creation of OTBI and BI Publisher reports.

TECHNICAL PROJECTS

AI Architect — AI-Native System Design Tool

Aug 2026 – Sep 2026

[GitHub](#) | [Live Demo](#)

- Built an AI tool that turns a plain-language description of a system into a real architecture diagram (services, databases, queues, infrastructure), complete with cost estimates and load testing, not just a picture generated by the AI.
- Built a cost and capacity model across ~30 component types with engine-specific pricing (CockroachDB vs. Postgres, DynamoDB vs. Redis), plus a deterministic load/failure simulator showing what breaks and in what order under a traffic spike or component failure.
- Built a repo-ingestion pipeline extracting source-attributed facts (routes, ORM models, migrations, API calls) from an imported GitHub repo, validating every proposed component against those citations; covered by 190+ backend tests.

AI Resume & Job Matcher — Full-Stack RAG Platform

Aug 2025 – Nov 2025 (Updated May 2026)

[GitHub](#) | [Live Demo](#)

- Built a full-stack AI recruitment platform (Next.js, FastAPI, Supabase, Qdrant) with a hybrid retrieval pipeline (dense + BM25, reciprocal rank fusion, cross-encoder reranking) grounding every AI verdict in a cited resume snippet.
- Designed a three-stage recruiter evaluation pipeline that shortlisted the top ~5% of candidates, reducing full rubric reviews by approximately 95%.
- Built an automated groundedness-evaluation harness (with OCR fallback for scanned resumes) and a corroboration layer that verifies weak resume claims against live GitHub data before trusting them.

Agentic RAG Customer Support Platform (Flipkart FAQ)

Nov 2025 – Jan 2026

[GitHub](#) | [Live Demo](#)

- Architected an end-to-end agentic RAG system using LlamaIndex (intent-routing planner, 5 tools including 4 MCP tools, LiteLLM gateway with fallback across 3 models/2 providers); LLM Guard guardrails cut latency from ~4.5s to ~1s, validated via 25-scenario RAGAS.
- Added vision-multimodal support (Groq Llama 4 Scout) so customers can attach photos of damaged items, grounding responses in real visual evidence.
- Built an intent-classification evaluation suite that caught a live production incident, two LLM models silently deprecated and failing through the fallback chain undetected, diagnosed, fixed, and verified against the live API.

EDUCATION

B.E. Computer Science and Engineering (Honours), SRM Easwari Engineering College (2022–2026) | CGPA: 8.8/10

CERTIFICATIONS & OPEN SOURCE

- Oracle Cloud Infrastructure – Agentic AI Certified Foundations Associate & AI Foundations Associate
- Machine Learning Specialization – Stanford University
- LangChain — Fixed a data-loss bug in RecursiveJsonSplitter; merged PR #39238 with 9 regression tests, 149 tests passing

TECHNICAL SKILLS

- Programming & Backend** — Python (Pandas, NumPy), FastAPI, SQL, MySQL, PostgreSQL (Supabase), Next.js, TypeScript
- LLM & Agentic AI** — LangChain, LangGraph, LlamaIndex, MCP, RAG Pipelines, Browser-Use Agents
- Vector Databases & Retrieval** — Vector Databases (Qdrant & Weaviate), Hybrid Search (BM25), Vision (Llama 4 Scout), OCR
- LLMOps & Evaluation** — LiteLLM, LLM Guard, Langfuse, RAGAS, OpenRouter, Groq API, OpenAI API
- Machine Learning** — Scikit-Learn, TensorFlow, NLP (NLTK & spaCy)
- Cloud & DevOps** — Git, GitHub, Docker, Nginx, Redis, Linux, Azure, Oracle Cloud Infrastructure (OCI)