

Soban Shankar

Chennai, India

+91-94443-89781 | ssoban2004@gmail.com | [Portfolio](#) | [GitHub](#) | [LinkedIn](#)

PROFESSIONAL SUMMARY

AI Engineer with hands-on enterprise AI experience through a 6-month internship and multiple end-to-end Generative AI projects. Skilled in building AI applications using LLMs, RAG, agentic AI systems, and MCP. Proficient in Python, LangChain, LangGraph, and vector databases.

EXPERIENCE

AI Engineering Intern | Drivestream | Chennai, India

Jan 2026 – Jul 2026

- Designed a multi-agent orchestration system to automate enterprise Oracle Fusion cloud transformation workflows, reducing manual configuration effort by approximately 60–70%.
- Developed a configuration automation agent that autonomously configured 20+ Oracle Fusion HCM and ERP modules, translating complex business rules into LLM-driven workflows.
- Built a conversational AI report agent using Retrieval-Augmented Generation (RAG) over 2,000+ Oracle Fusion HCM table and view definitions, enabling schema-aware SQL generation and natural language creation of OTBI and BI Publisher reports.

TECHNICAL PROJECTS

AI Resume & Job Matcher — Full-Stack RAG Platform

Aug 2025 – Nov 2025 (Updated May 2026)

[GitHub](#) | [Live Demo](#)

- Built a full-stack AI recruitment platform using Next.js 16, FastAPI, Supabase, and Qdrant, enabling job seekers to receive AI-powered resume analysis and optimization while recruiters manage end-to-end hiring workflows.
- Engineered a hybrid retrieval pipeline combining dense and sparse (BM25) search with reciprocal rank fusion and cross-encoder reranking, enabling grounded AI responses with citation-backed resume evidence.
- Designed a three-stage recruiter evaluation pipeline (dense prescreen → skill match → LLM review) that shortlisted the top ~5% of candidates, reducing full rubric reviews by approximately 95%.
- Implemented a three-tier LLM fallback chain (Ollama → Gemini → Groq) to improve resilience against rate limits, with Supabase authentication and a PostgreSQL backend supporting persistent, reopenable hiring campaigns.
- Added a corroboration layer that verifies weak resume claims against live GitHub repository content before upgrading their evidence score, closing a keyword-stuffing gap that flat keyword/LLM matchers don't address.

Agentic RAG Customer Support Platform (Flipkart FAQ)

Nov 2025 – Jan 2026

[GitHub](#) | [Live Demo](#)

- Built an agentic RAG customer support platform over a ~2,000-entry Flipkart FAQ corpus, enabling automated customer support with order status, refund tracking, and human escalation workflows through Model Context Protocol (MCP) tools.
- Architected an end-to-end agentic RAG system using LlamaIndex, featuring an intent-routing planner, 5 integrated tools (FAQ search + 4 MCP tools), and a LiteLLM gateway with automatic fallback across 3 models and 2 providers.
- Implemented LLM Guard guardrails with 7 parallel scanners, reducing latency from approximately 4.5s to 1s, validated through a 25-scenario RAGAS evaluation combining hybrid retrieval, reranking, and Qdrant vector search, and deployed as a Docker container on Render.

EDUCATION

B.E. Computer Science and Engineering (Honours), SRM Easwari Engineering College (2022–2026) | CGPA: 8.8/10

CERTIFICATIONS & OPEN SOURCE

- Oracle Cloud Infrastructure – Agentic AI Certified Foundations Associate
- Oracle Cloud Infrastructure – AI Foundations Associate
- Machine Learning Specialization – Stanford University
- LangChain — Fixed a data-loss bug in RecursiveJsonSplitter (Issue #39192); implemented and merged PR #39238 with 9 regression tests, validated against 149 passing tests

TECHNICAL SKILLS

- Programming & Backend** — Python (Pandas, NumPy), FastAPI, SQL, MySQL, PostgreSQL (Supabase)
- Frontend** — Next.js, TypeScript
- LLM & Agentic AI** — LangChain, LangGraph, LlamaIndex, MCP (Model Context Protocol), RAG Pipelines, Prompt Engineering
- Speech & Voice AI** — Deepgram Nova-3 (STT), Speaker Diarization, Multilingual/Code-Switched Audio
- Vector Databases & Retrieval** — Vector Databases (Qdrant & Weaviate), Hybrid Search (BM25) & Reranking, Hugging Face Models
- LLMOps & Evaluation** — LiteLLM, LLM Guard, Langfuse, RAGAS, OpenRouter, Groq API, OpenAI API
- Machine Learning** — Scikit-Learn, TensorFlow, NLP (NLTK & spaCy)
- Cloud & DevOps** — Git, GitHub, Docker, Nginx, Redis, Linux, Azure, Oracle Cloud Infrastructure (OCI)